

8.2 Red-Black Trees

A *red-black tree* is a binary search tree with one extra attribute for each node: the *colour*, which is either red or black. We also need to keep track of the parent of each node, so that a red-black tree's node structure would be:


```
struct t_red_black_node {
    enum { red, black } colour;
    void *item;
    struct t_red_black_node *left,
                             *right,
                             *parent;
}
```

For the purpose of this discussion, the NULL nodes which terminate the tree are considered to be the leaves and are coloured black.


Definition of a red-black tree

A red-black tree is a binary search tree which has the following *red-black properties*:

1. Every node is either red or black.
 2. Every leaf (NULL) is black.
 3. If a node is red, then both its children are black.
 4. Every simple path from a node to a descendant leaf contains the same number of black nodes.
3. implies that on any path from the root to a leaf, red nodes must not be adjacent. However, any number of black nodes may appear in a sequence.



A basic red-black tree



Basic red-black tree with the **sentinel** nodes added. Implementations of the red-black tree algorithms will usually include the sentinel nodes as a convenient means of flagging that you have reached a leaf node.

They are the NULL black nodes of property 2.

The number of black nodes on any path from, but not including, a node x to a leaf is called the *black-height* of a node, denoted $\mathbf{bh}(x)$. We can prove the following lemma:

Lemma

A red-black tree with n internal nodes has height at most $2\log(n+1)$.
(For a proof, see Cormen, p 264)

This demonstrates why the red-black tree is a good search tree: it can always be searched in $O(\log n)$ time.


As with heaps, additions and deletions from red-black trees destroy the red-black property, so we need to restore it. To do this we need to look at some operations on red-black trees.

Rotations

A rotation is a local operation in a search tree that preserves *in-order* traversal key ordering.

Note that in both trees, an in-order traversal yields:

A x B y C



The left_rotate operation may be encoded:

```
left_rotate( Tree T, node x ) {
    node y;
    y = x->right;
    /* Turn y's left sub-tree into x's right sub-tree */
    x->right = y->left;
    if ( y->left != NULL )
        y->left->parent = x;
    /* y's new parent was x's parent */
    y->parent = x->parent;
    /* Set the parent to point to y instead of x */
    /* First see whether we're at the root */
    if ( x->parent == NULL ) T->root = y;
    else
        if ( x == (x->parent)->left )
            /* x was on the left of its parent */
            x->parent->left = y;
        else
            /* x must have been on the right */
            x->parent->right = y;
    /* Finally, put x on y's left */
    y->left = x;
    x->parent = y;
}
```

Insertion

Insertion is somewhat complex and involves a number of cases. Note that we start by inserting the new node, x , in the tree just as we would for any other binary tree, using the `tree_insert` function. This new node is labelled red, and possibly destroys the red-black property. The main loop moves up the tree, restoring the red-black property.

```
rb_insert( Tree T, node x ) {
    /* Insert in the tree in the usual way */
    tree_insert( T, x );
    /* Now restore the red-black property */
    x->colour = red;
    while ( (x != T->root) && (x->parent->colour == red) ) {
        if ( x->parent == x->parent->parent->left ) {
            /* If x's parent is a left, y is x's right 'uncle' */
            y = x->parent->parent->right;
            if ( y->colour == red ) {
                /* case 1 - change the colours */
                x->parent->colour = black;
                y->colour = black;
                x->parent->parent->colour = red;
            }
        }
    }
}
```

```

        /* Move x up the tree */
        x = x->parent->parent;
    }
else {
    /* y is a black node */
    if ( x == x->parent->right ) {
        /* and x is to the right */
        /* case 2 - move x up and rotate */
        x = x->parent;
        left_rotate( T, x );
    }
    /* case 3 */
    x->parent->colour = black;
    x->parent->parent->colour = red;
    right_rotate( T, x->parent->parent );
}
}
else {
    /* repeat the "if" part with right and left
    exchanged */
}
}
/* Colour the root black */
T->root->colour = black;
}

```

[Here's an example of the insertion operation.](#)

Animation

Red-Black Tree Animation

This animation was written by Linda Luo, Mervyn Ng, Anita Lee, John Morris and Woi Ang.

Please email comments to:

John Morris j.morris@auckland.ac.nz

Examination of the code reveals only one loop. In that loop, the node at the root of the sub-tree whose red-black property we are trying to restore, x , may be moved up the tree *at least one level* in each iteration of the loop. Since the tree originally has $O(\log n)$ height, there are $O(\log n)$ iterations. The `tree_insert` routine also has $O(\log n)$ complexity, so overall the `rb_insert` routine also has $O(\log n)$ complexity.

Key terms

Red-black trees

Trees which remain **balanced** - and thus guarantee $O(\log n)$ search times - in a dynamic environment. Or more importantly, since any tree can be re-balanced - but at considerable cost - can be re-balanced in $O(\log n)$ time.

Continue on to [AVL Trees](#)

Back to the [Table of Contents](#)